

Connected to Python 3.12.10

```
In [ ]: # %pip install ucimlrepo
from ucimlrepo import fetch_ucirepo

# fetch dataset
us_census_data_1990 = fetch_ucirepo(id=116)

# This project uses the 1990 U.S. Census dataset from the UC
```

```
In [ ]: # Load data (as pandas dataframes)
X = us_census_data_1990.data.features
y = us_census_data_1990.data.targets

# Combine features and targets (if targets exist)
df = us_census_data_1990.data.original

print(df.shape)           # rows, columns
print(df.head())          # peek at first 5 rows
```

(1534490, 69)

	caseid	dAge	dAncstry1	dAncstry2	iAvail	iCitizen	iCla
ss	dDepart	\					
0	10000	5	0	1	0	0	0
5		3					
1	10001	6	1	1	0	0	0
7		5					
2	10002	3	1	2	0	0	0
7		4					
3	10003	4	1	2	0	0	0
1		3					
4	10004	7	1	1	0	0	0
0		0					

	iDisabl1	iDisabl2	...	iTmpabsnt	dTravtime	iVietnam	d
Week89	iWork89	\					
0	2	2	...	0.0	5.0	0.0	0.0
2.0	1.0						
1	2	2	...	0.0	1.0	0.0	0.0
2.0	1.0						
2	2	2	...	0.0	2.0	0.0	0.0
2.0	1.0						
3	2	2	...	0.0	1.0	0.0	0.0
1.0	1.0						
4	2	2	...	3.0	0.0	0.0	0.0
0.0	2.0						

	iWorklwk	iWWII	iYearsch	iYearwrk	dYrsserv
0	1.0	0.0	11.0	1.0	0.0
1	1.0	0.0	5.0	1.0	0.0
2	1.0	0.0	10.0	1.0	0.0
3	1.0	0.0	10.0	1.0	0.0
4	2.0	0.0	5.0	6.0	0.0

[5 rows x 69 columns]

```
In [ ]: # Exploratory Data Analysis (EDA)
# Check dimensions
print(X.shape)
```

```
import seaborn as sns
import matplotlib.pyplot as plt
import pandas as pd
```

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report

# Step 1: Separate features and target
X = df.drop(columns=['dTravtime'])
y = df['dTravtime'].astype('category') # categorical outcome

# Step 2: Mark categorical predictors
categorical_vars = [
    "iSex", "dOccup", "dIndustry", "dAncstry1", "dAncstry2",
    "dHispanic", "dPOB", "iCitizen", "iEnglish", "iLang1",
    "iMilitary", "iSchool", "iWorklwk", "iDisab11",
    "iDisab12", "iMobillim", "iVietnam", "iWWII"
]
numeric_vars = [col for col in X.columns if col not in categorical_vars]

# Define numeric and categorical variables from top features
numeric_features = [
    "dDepart", "iMeans", "dHours", "iClass", "dHour89", "iRi",
    "dWeek89", "dRearning", "iYearsch", "iYearwrk", "dIncome",
    "iLooking", "dAge", "dPwgt1", "iRPOB", "iTmpabsnt"
]

categorical_features = [
    "dOccup", "dIndustry", "iWorklwk"
]

# Function to plot numeric features
def plot_numeric_univariate(df, features, bins=30):
    for col in features:
        if col in df.columns:
            plt.figure(figsize=(6,4))
            sns.histplot(df[col], bins=bins, kde=False, color='blue')
            plt.title(f"Distribution of {col}")
            plt.xlabel(col)
```

```
plt.ylabel("Count")
plt.grid(True, alpha=0.3)
plt.show()

# Function to plot categorical features
def plot_categorical_univariate(df, features, top_n=10):
    for col in features:
        if col in df.columns:
            plt.figure(figsize=(8,4))
            counts = df[col].value_counts().nlargest(top_n)
            sns.barplot(x=counts.index.astype(str), y=counts)
            plt.title(f"Top {top_n} Categories of {col}")
            plt.xlabel(col)
            plt.ylabel("Count")
            plt.xticks(rotation=45)
            plt.grid(True, alpha=0.3)
            plt.show()

# Run univariate plots on your sample dataframe
plot_numeric_univariate(df, numeric_features)
plot_categorical_univariate(df, categorical_features)

# Bivariate EDA
# Distribution of commute time categories
sns.countplot(x=y)
plt.xlabel("Commute time category (0-6)")
plt.ylabel("Count")
plt.title("Distribution of Commute Time Categories")
plt.show()

# Compare commute time categories vs age and earnings
sns.boxplot(x=y, y=df['dAge'])
plt.title("Commute Time Category vs Age")
plt.show()

sns.boxplot(x=y, y=df['dRearning'])
plt.title("Commute Time Category vs Earnings")
plt.show()

from scipy.stats import f_oneway, chi2_contingency

# Target variable
```

```

y = df['dTravtime'].astype('category')

# Separate numeric and categorical predictors
numeric_features = [
    "dDepart", "iMeans", "dHours", "iClass", "dHour89", "iRiders",
    "dWeek89", "dRearning", "iYearsch", "iYearwrk", "dIncome1",
    "iLooking", "dAge", "dPwgt1", "iRPOB", "iTmptabsnt"
]

categorical_features = ["dOccup", "dIndustry", "iWorklwk"]

# ANOVA for numeric predictors vs. commute time
anova_results = {}
for col in numeric_features:
    if col in df.columns:
        groups = [df[col][y==cat] for cat in y.cat.categories]
        try:
            fstat, pval = f_oneway(*groups)
            anova_results[col] = pval
        except Exception:
            anova_results[col] = np.nan

print("ANOVA p-values (numeric predictors vs. commute time):")
print(pd.Series(anova_results).sort_values())

# Chi-square for categorical predictors vs. commute time
chi2_results = {}
for col in categorical_features:
    if col in df.columns:
        table = pd.crosstab(df[col].astype(str), y.astype(str))
        try:
            chi2, pval, _, _ = chi2_contingency(table)
            chi2_results[col] = pval
        except Exception:
            chi2_results[col] = np.nan

print("\nChi-square p-values (categorical predictors vs. commute time):")
print(pd.Series(chi2_results).sort_values())

import seaborn as sns
import matplotlib.pyplot as plt


# Compute correlation matrix for numeric features

```


```
numeric_df = df[numeric_features].dropna()
corr_matrix = numeric_df.corr(method="pearson")

# Plot heatmap
plt.figure(figsize=(10,8))
sns.heatmap(corr_matrix, annot=True, cmap="coolwarm", center=0)
plt.title("Pearson Correlation Heatmap (Numeric Predictors)")
plt.show()
```


(1534490, 68)




Distribution of iMeans




Distribution of dHours




Distribution of iClass




Distribution of dHour89




Distribution of iRiders




Distribution of dWeek89




Distribution of dRearning




Distribution of iYearsch




Distribution of iYearwrk




Distribution of dIncome1







Distribution of dPwgt1



Distribution of iRPOB






```
<ipython-input-3-3cd4f2b3c74b>:62: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.


```
    sns.barplot(x=counts.index.astype(str), y=counts.values, palette="viridis")
```



```
<ipython-input-3-3cd4f2b3c74b>:62: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
    sns.barplot(x=counts.index.astype(str), y=counts.values, palette="viridis")
```




```
<ipython-input-3-3cd4f2b3c74b>:62: FutureWarning:
```


Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
    sns.barplot(x=counts.index.astype(str), y=counts.values, palette="viridis")
```


Top 10 Categories of iWorklwk




Distribution of Commute Time Categories



Commute Time Category vs Age



Commute Time Category vs Earnings




ANOVA p-values (numeric predictors vs. commute time):

```
dDepart      0.0
iMeans       0.0
dHours       0.0
iClass        0.0
dHour89      0.0
iRiders       0.0
dWeek89      0.0
dRearing      0.0
iYearsch     0.0
iYearwrk     0.0
dIncome1      0.0
iLooking      0.0
dAge          0.0
dPwgt1        0.0
iRPOB         0.0
iTmpabsnt    0.0
dtype: float64
```

Chi-square p-values (categorical predictors vs. commute time):

```
dOccup       0.0
dIndustry     0.0
iWorklwk     0.0
dtype: float64
```



```
In [ ]: # Multinomial Logistic Regression with Backward Stepwise Selection
import warnings
warnings.filterwarnings("ignore")

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report
from scipy.stats import f_oneway, chi2_contingency
from collections import Counter

# Step 1: Sample 50,000 rows
df_sample = df.sample(n=50000, random_state=42)

# Step 2: Collapse high-cardinality categorical variables
def collapse_categories(series, top_n=4):
    top_cats = series.value_counts().nlargest(top_n).index
```

```

        return series.where(series.isin(top_cats), other='Other')

for col in ['dIndustry','dOccup','dAncstry1','dAncstry2','il
    if col in df_sample.columns:
        df_sample[col] = collapse_categories(df_sample[col],

# Ensure categorical columns are all strings
for col in ['dIndustry','dOccup','dAncstry1','dAncstry2','il
    if col in df_sample.columns:
        df_sample[col] = collapse_categories(df_sample[col],

# Also cast other categorical predictors to string
categorical_vars = [
    "iSex", "dOccup", "dIndustry", "dAncstry1", "dAncstry2",
    "dHispanic", "dPOB", "iCitizen", "iEnglish", "iLang1",
    "iMilitary", "iSchool", "iWorklwk", "iDisabl1",
    "iDisabl2", "iMobillim", "iVietnam", "iWWII"
]
for col in categorical_vars:
    if col in df_sample.columns:
        df_sample[col] = df_sample[col].astype(str)

# Step 3: Target and predictors
y = df_sample['dTravtime'].astype('category')
predictors = [col for col in df_sample.columns if col not in

# Step 4: Pre-screen predictors
scores = {}
for col in predictors:
    if pd.api.types.is_numeric_dtype(df_sample[col]):
        groups = [df_sample[col][y==cat] for cat in y.cat.categories]
        try:
            _, pval = f_oneway(*groups)
            scores[col] = pval
        except Exception:
            scores[col] = 1.0
    else:
        table = pd.crosstab(df_sample[col].astype(str), y.astype(str))
        try:
            _, pval, _, _ = chi2_contingency(table)
            scores[col] = pval
        except Exception:
            scores[col] = 1.0

```

```
# Step 5: Select top 20 predictors
top_predictors = sorted(scores, key=scores.get)[:20]
print("Top 20 predictors:", top_predictors)

# Step 6: Build design matrix using top predictors
X = df_sample[top_predictors]
y_codes = y.cat.codes

# Step 7: Balance classes (undersample majority)
counts = Counter(y_codes)
min_count = min(counts.values())
balanced_idx = np.hstack([
    np.random.choice(np.where(y_codes == cls)[0], min_count,
                     replace=False) for cls in counts.keys()])
X_balanced = X.iloc[balanced_idx]
y_balanced = y_codes.iloc[balanced_idx]

print("Balanced class counts:", Counter(y_balanced))

# Step 8: Train/test split
X_train, X_test, y_train, y_test = train_test_split(
    X_balanced, y_balanced, test_size=0.2, random_state=42,
)

# Step 9: Build preprocessing pipeline (restricted to available columns)
def make_pipeline(predictor_list):
    categorical_vars = [col for col in predictor_list if col in [
        "iSex", "dOccup", "dIndustry", "dAncstry1", "dAncstr",
        "dHispanic", "dPOB", "iCitizen", "iEnglish", "iLang1",
        "iMilitary", "iSchool", "iWorklwk", "iDisabl1",
        "iDisabl2", "iMobillim", "iVietnam", "iWWII"
    ]]
    numeric_vars = [col for col in predictor_list if col not in categorical_vars]

    from sklearn.preprocessing import StandardScaler, OneHotEncoder
    from sklearn.compose import ColumnTransformer
    from sklearn.pipeline import Pipeline

    preprocessor = ColumnTransformer(
        transformers=[
            ("num", StandardScaler(), numeric_vars),
            ("cat", OneHotEncoder(), categorical_vars)
        ]
    )
    return Pipeline([
        ("preprocessor", preprocessor),
        ("logit", LogisticRegression())
    ])

```

```
( "cat", OneHotEncoder(handle_unknown="ignore"),
    )
)

pipeline = Pipeline(steps=[
    ("preprocess", preprocess),
    ("model", LogisticRegression(
        multi_class="multinomial", solver="lbfgs", pena]
    )))
return pipeline

# Initial model
logit_pipeline = make_pipeline(list(X_train.columns))
logit_pipeline.fit(X_train, y_train)
init_acc = accuracy_score(y_test, logit_pipeline.predict(X_t
print("Initial accuracy with 20 predictors (balanced):", ini

# Step 10: Backward stepwise selection
def backward_stepwise(X_train, y_train, X_test, y_test, thre
    included = list(X_train.columns)
    acc_history = []
    pipeline = make_pipeline(included)
    pipeline.fit(X_train, y_train)
    best_acc = accuracy_score(y_test, pipeline.predict(X_te
    acc_history.append((len(included), best_acc, included.co

    improved = True
    while improved and len(included) > min_vars:
        improved = False
        for col in included:
            trial = [c for c in included if c != col]
            trial_pipeline = make_pipeline(trial)
            trial_pipeline.fit(X_train[trial], y_train)
            acc = accuracy_score(y_test, trial_pipeline.predict(X_t
            if acc >= best_acc - threshold:
                included.remove(col)
                best_acc = acc
                acc_history.append((len(included), acc, inclu
                print(f"Dropped {col}, accuracy {acc}")
                improved = True
                break
    return acc_history
```

```

acc_history = backward_stepwise(X_train, y_train, X_test, y_
                                 _test)

# Step 11: Plot accuracy vs. predictors retained
steps, accs, _ = zip(*acc_history)
plt.figure(figsize=(8,5))
plt.plot(steps, accs, marker='o')
plt.gca().invert_xaxis()
plt.xlabel("Number of Predictors Retained")
plt.ylabel("Test Accuracy")
plt.title("Backward Stepwise Selection (Balanced, 50k): Accu-
racy vs. Number of Predictors Retained")
plt.grid(True)
plt.show()

# Step 12: Extract best model
best_step = max(acc_history, key=lambda x: x[1])
best_vars = best_step[2]
print("Best model predictors:", best_vars)
print("Best model accuracy:", best_step[1])

best_pipeline = make_pipeline(best_vars)
best_pipeline.fit(X_train[best_vars], y_train)

# Coefficient table
best_model = best_pipeline.named_steps["model"]
coef_df = pd.DataFrame(best_model.coef_, columns=best_pipeline.
                        get_params()["model"].estimator_.classes_)
coef_df.index = [f"Class {c}" for c in best_model.classes_]
print("Coefficient table for best model (balanced, 50k):")
print(coef_df)

# ROC curves for best model
from sklearn.preprocessing import label_binarize
from sklearn.metrics import roc_curve, auc

y_test_bin = label_binarize(y_test, classes=np.unique(y_test))
y_score = best_pipeline.predict_proba(X_test[best_vars])

fpr, tpr, roc_auc = {}, {}, {}
plt.figure(figsize=(8,6))
for i in range(y_test_bin.shape[1]):
    fpr[i], tpr[i], _ = roc_curve(y_test_bin[:, i], y_score[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])
    plt.plot(fpr[i], tpr[i], label=f"Class {i} (AUC = {roc_a

```


```
plt.plot([0,1],[0,1],'k--')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("Multiclass ROC (OvR) for Best Multinomial Logisti
plt.legend()
plt.show()

# We get between 32 and 39% accuracy with 12 predictors rete

# Multinomial logistic regression has served us well as a tr
# That said, the approach has limitations. It assumes linear
# In short, multinomial logit is a strong pedagogical and di
```

Top 20 predictors: ['dAge', 'iClass', 'dDepart', 'iDisabl1',
'iDisabl2', 'dHour89', 'dHours', 'dIncome1', 'dIncome5', 'dIn
dustry', 'iLang1', 'iLooking', 'iMarital', 'iMeans', 'iMilita
ry', 'iMobility', 'iMobillim', 'dOccup', 'iPerscare', 'dPover
ty']
Balanced class counts: Counter({5: 1331, 0: 1331, 2: 1331, 4:
1331, 3: 1331, 1: 1331, 6: 1331})
Initial accuracy with 20 predictors (balanced): 0.34281115879
82833
Dropped iClass, accuracy 0.3449570815450644
Dropped iDisabl1, accuracy 0.3449570815450644
Dropped dAge, accuracy 0.34549356223175964
Dropped iDisabl2, accuracy 0.34549356223175964
Dropped dHour89, accuracy 0.34603004291845496
Dropped dHours, accuracy 0.34549356223175964
Dropped dIncome5, accuracy 0.34763948497854075
Dropped iLang1, accuracy 0.35354077253218885
Dropped iLooking, accuracy 0.35354077253218885
Dropped iMobility, accuracy 0.3540772532188841
Dropped iMarital, accuracy 0.3546137339055794
Dropped iMobillim, accuracy 0.3562231759656652

Backward Stepwise Selection (Balanced, 50k): Accuracy vs. Predictors



Best model predictors: ['dDepart', 'dIncome1', 'dIndustry', 'iMeans', 'iMilitary', 'dOccup', 'iPerscare', 'dPoverty']
 Best model accuracy: 0.3562231759656652

Coefficient table for best model (balanced, 50k):

	num_dDepart	num_dIncome1	num_iMeans	num_iPers
care \				
Class 0	-8.571311	-0.730035	-0.088052	-0.15
3620				
Class 1	1.666998	-0.107146	0.220191	0.17
9954				
Class 2	1.539204	0.007208	0.081969	0.11
7289				
Class 3	1.621859	0.065694	-0.119641	-0.14
0720				
Class 4	1.557053	0.168993	-0.197709	-0.08
2571				
Class 5	1.267848	0.284499	-0.077638	0.07
6030				
Class 6	0.918350	0.310788	0.180880	0.00
3639				

	num_dPoverty	cat_dIndustry_0.0	cat_dIndustry_7.
0 \			
Class 0	-0.048645	0.584788	-0.43495
6			
Class 1	-0.080234	-0.066610	0.13678
7			
Class 2	0.012837	-0.087874	0.21537
4			
Class 3	0.008620	-0.104121	0.10281
4			
Class 4	0.094931	-0.090354	0.03589
4			
Class 5	0.034221	-0.092932	0.00800
5			
Class 6	-0.021730	-0.142897	-0.06391
8			

	cat_dIndustry_9.0	cat_dIndustry_Other	cat_iMili
aty_0.0 \			
Class 0	-0.101257	-0.415726	
0.196323			
Class 1	0.131661	-0.136821	-

0.014312			
Class 2	0.073186	-0.131235	-
0.023320			
Class 3	0.021155	0.040706	-
0.054329			
Class 4	0.002019	0.114746	-
0.039497			
Class 5	-0.104756	0.252170	-
0.023635			
Class 6	-0.022007	0.276160	-
0.041230			


	cat__iMilitary_1.0	cat__iMilitary_2.0	cat__iMilita
ry_3.0 \			
Class 0	-0.243171	-0.223810	0.
091776			
Class 1	0.142776	0.173319	-0.
173485			
Class 2	0.467222	-0.110083	-0.
226605			
Class 3	0.491270	-0.142478	-0.
135350			
Class 4	0.080830	-0.062747	0.
072730			
Class 5	-0.210814	0.080470	0.
052652			
Class 6	-0.728114	0.285328	0.
318282			

	cat__iMilitary_4.0	cat__d0ccup_0.0	cat__d0ccup_1.0
\			
Class 0	-0.188269	0.584788	-0.184130
Class 1	-0.063282	-0.066610	0.041648
Class 2	-0.037763	-0.087874	0.064463
Class 3	-0.098560	-0.104121	-0.027677
Class 4	0.010988	-0.090354	0.041825
Class 5	0.163815	-0.092932	0.096909
Class 6	0.213072	-0.142897	-0.033038

	cat__d0ccup_2.0	cat__d0ccup_Other
Class 0	-0.087370	-0.680439
Class 1	-0.057498	0.147477
Class 2	0.007819	0.085044

Class 3	0.040121	0.152230
Class 4	0.035333	0.075501
Class 5	-0.038769	0.097278
Class 6	0.100364	0.122909

Multiclass ROC (OvR) for Best Multinomial Logistic Regression



```
In [ ]: # K-Nearest Neighbors (KNN) with PCA
import warnings
warnings.filterwarnings("ignore")

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, f1_score, classification_report
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.decomposition import PCA
from sklearn.neighbors import KNeighborsClassifier
from collections import Counter

# Step 1: Sample 50,000 rows
```

```

df_sample = df.sample(n=50000, random_state=42)

# Step 2: Collapse high-cardinality categorical variables
def collapse_categories(series, top_n=4):
    top_cats = series.value_counts().nlargest(top_n).index
    return series.where(series.isin(top_cats), other='Other')

for col in ['dIndustry', 'dOccup', 'dAncstry1', 'dAncstry2', 'il':
    if col in df_sample.columns:
        df_sample[col] = collapse_categories(df_sample[col], top_n=4)

# Cast other categorical predictors to string
categorical_vars_master = [
    "iSex", "dOccup", "dIndustry", "dAncstry1", "dAncstry2",
    "dHispanic", "dPOB", "iCitizen", "iEnglish", "iLang1",
    "iMilitary", "iSchool", "iWorklwk", "iDisabl1",
    "iDisabl2", "iMobillim", "iVietnam", "iWWII"
]
for col in categorical_vars_master:
    if col in df_sample.columns:
        df_sample[col] = df_sample[col].astype(str)

# Step 3: Target and predictors
y = df_sample['dTravtime'].astype('category')
predictors = [col for col in df_sample.columns if col not in ['dTravtime']]
X = df_sample[predictors]
y_codes = y.cat.codes

# Step 4: Balance classes (undersample majority)
counts = Counter(y_codes)
min_count = min(counts.values())
balanced_idx = np.hstack([
    np.random.choice(np.where(y_codes == cls)[0], min_count,
                    replace=False) for cls in counts.keys()])
X_balanced = X.iloc[balanced_idx]
y_balanced = y_codes.iloc[balanced_idx]

print("Balanced class counts:", Counter(y_balanced))

# Step 5: Train/test split
X_train, X_test, y_train, y_test = train_test_split(
    X_balanced, y_balanced, test_size=0.2, random_state=42,
    stratify=y_balanced)

```

```
)  
  
# Step 6: Build preprocessing pipeline  
def make_pipeline_knn(predictor_list, n_neighbors=5):  
    categorical_vars = [col for col in predictor_list if col]  
    numeric_vars = [col for col in predictor_list if col not in categorical_vars]  
  
    preprocessor = ColumnTransformer(  
        transformers=[  
            ("num", StandardScaler(), numeric_vars),  
            ("cat", OneHotEncoder(handle_unknown="ignore")),  
        ]  
    )  
  
    pipeline = Pipeline(steps=[  
        ("preprocess", preprocessor),  
        ("pca", PCA(n_components=0.95, random_state=42)), #  
        ("model", KNeighborsClassifier(n_neighbors=n_neighbors))  
    ])  
    return pipeline  
  
# Step 7: Hyperparameter tuning for n_neighbors  
results = []  
best_model = None  
best_f1 = -1  
  
for k in [3, 5, 7, 9]:  
    knn_pipeline = make_pipeline_knn(list(X_train.columns),  
    knn_pipeline.fit(X_train, y_train)  
    y_pred = knn_pipeline.predict(X_test)  
    acc = accuracy_score(y_test, y_pred)  
    macro_f1 = f1_score(y_test, y_pred, average="macro")  
    results.append((k, acc, macro_f1))  
    print(f"\nKNN (k={k})")  
    print("Test accuracy:", acc)  
    print("Macro F1:", macro_f1)  
    print("Classification report:")  
    print(classification_report(y_test, y_pred))  
    if macro_f1 > best_f1:  
        best_f1 = macro_f1  
        best_model = knn_pipeline  
  
# Step 8: ROC curves for best model
```

```
y_test_bin = label_binarize(y_test, classes=np.unique(y_test)
y_score = best_model.predict_proba(X_test)

fpr, tpr, roc_auc = {}, {}, {}
plt.figure(figsize=(8,6))
for i in range(y_test_bin.shape[1]):
    fpr[i], tpr[i], _ = roc_curve(y_test_bin[:, i], y_score[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])
    plt.plot(fpr[i], tpr[i], label=f"Class {i} (AUC = {roc_auc[i]:.2f})")

plt.plot([0,1],[0,1],'k--')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("Multiclass ROC (OvR) for Best Balanced KNN Model")
plt.legend()
plt.show()

# Our PCA cumulative explained variance plot shows that near
# We see that k=5 gives the best test accuracy of around 30%
```

Balanced class counts: Counter({5: 1331, 0: 1331, 2: 1331, 4: 1331, 3: 1331, 1: 1331, 6: 1331})

KNN (k=3)

Test accuracy: 0.3084763948497854

Macro F1: 0.311146807929259

Classification report:

	precision	recall	f1-score	support
0	1.00	0.95	0.98	266
1	0.22	0.23	0.23	266
2	0.16	0.18	0.17	266
3	0.22	0.21	0.21	267
4	0.19	0.18	0.19	266
5	0.19	0.19	0.19	266
6	0.23	0.21	0.22	267
accuracy			0.31	1864
macro avg	0.31	0.31	0.31	1864
weighted avg	0.31	0.31	0.31	1864

KNN (k=5)

Test accuracy: 0.3084763948497854

Macro F1: 0.3112538395355052

Classification report:

	precision	recall	f1-score	support
0	0.99	0.95	0.97	266
1	0.23	0.24	0.23	266
2	0.17	0.20	0.18	266
3	0.18	0.18	0.18	267
4	0.18	0.18	0.18	266
5	0.21	0.22	0.22	266
6	0.23	0.20	0.22	267
accuracy			0.31	1864
macro avg	0.31	0.31	0.31	1864
weighted avg	0.31	0.31	0.31	1864

KNN (k=7)

Test accuracy: 0.3133047210300429

Macro F1: 0.3151779927794175

Classification report:

	precision	recall	f1-score	support
0	0.99	0.95	0.97	266
1	0.25	0.26	0.26	266
2	0.20	0.23	0.21	266
3	0.19	0.18	0.18	267
4	0.17	0.18	0.18	266
5	0.21	0.21	0.21	266
6	0.21	0.19	0.20	267
accuracy			0.31	1864
macro avg	0.32	0.31	0.32	1864
weighted avg	0.32	0.31	0.32	1864

KNN (k=9)


Test accuracy: 0.3197424892703863

Macro F1: 0.3214678334062072

Classification report:

	precision	recall	f1-score	support
0	0.99	0.95	0.97	266
1	0.27	0.29	0.28	266
2	0.18	0.20	0.19	266
3	0.19	0.19	0.19	267
4	0.18	0.17	0.17	266
5	0.20	0.20	0.20	266
6	0.25	0.23	0.24	267
accuracy			0.32	1864
macro avg	0.32	0.32	0.32	1864
weighted avg	0.32	0.32	0.32	1864

Multiclass ROC (OvR) for Best Balanced KNN Model



```
In [ ]: # Single Tree, pruned
import warnings
warnings.filterwarnings("ignore")

import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import label_binarize
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, classification_report
import matplotlib.pyplot as plt
from collections import Counter

# Step 1: Sample 100,000 rows
df_sample = df.sample(n=100000, random_state=42)

# Step 2: Collapse high-cardinality categorical variables
def collapse_categories(series, top_n=4):
    top_cats = series.value_counts().nlargest(top_n).index
    return series.where(series.isin(top_cats), other='Other')
```

```
for col in ['dIndustry', 'dOccup', 'dAncstry1', 'dAncstry2', 'il':
    if col in df_sample.columns:
        df_sample[col] = collapse_categories(df_sample[col],)

# Step 3: Target and predictors
y = df_sample['dTravtime'].astype('category')
predictors = [col for col in df_sample.columns if col not in y]

# Step 4: One-hot encode
X = pd.get_dummies(df_sample[predictors], drop_first=True)
y_codes = y.cat.codes

# Step 5: Balance classes (undersample majority)
counts = Counter(y_codes)
min_count = min(counts.values())
balanced_idx = np.hstack([
    np.random.choice(np.where(y_codes==cls)[0], min_count, replace=False)
    for cls in counts.keys()
])
X_balanced = X.iloc[balanced_idx]
y_balanced = y_codes.iloc[balanced_idx]

print("Balanced class counts:", Counter(y_balanced))

# Step 6: Train/test split
X_train, X_test, y_train, y_test = train_test_split(
    X_balanced, y_balanced, test_size=0.2, random_state=42,
)

# Step 7: Grid search for pruning
param_grid = {
    "max_depth": [5, 10, 15],
    "min_samples_leaf": [5, 10, 20],
    "min_samples_split": [10, 20]
}

grid = GridSearchCV(
    DecisionTreeClassifier(random_state=42),
    param_grid,
    cv=3,
    scoring="f1_macro"
)
grid.fit(X_train, y_train)
```

```
best_tree = grid.best_estimator_
print("Best parameters:", grid.best_params_)
print("Best macro F1 (CV):", grid.best_score_)

# Step 8: Evaluate best tree
y_pred = best_tree.predict(X_test)
acc = accuracy_score(y_test, y_pred)
print("Decision Tree test accuracy:", acc)
print("Classification report:")
print(classification_report(y_test, y_pred))


# Step 9: ROC curves (OvR)
y_test_bin = label_binarize(y_test, classes=np.unique(y_test))
y_score = best_tree.predict_proba(X_test)

fpr, tpr, roc_auc = {}, {}, {}
plt.figure(figsize=(8,6))
for i in range(y_test_bin.shape[1]):
    fpr[i], tpr[i], _ = roc_curve(y_test_bin[:, i], y_score[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])
    plt.plot(fpr[i], tpr[i], label=f"Class {i} (AUC = {roc_auc[i]:.2f})")

plt.plot([0,1],[0,1],'k--')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("Multiclass ROC (OvR) for Tuned Decision Tree")
plt.legend()
plt.show()
```

Balanced class counts: Counter({5: 2692, 0: 2692, 2: 2692, 4: 2692, 3: 2692, 1: 2692, 6: 2692})
 Best parameters: {'max_depth': 10, 'min_samples_leaf': 20, 'min_samples_split': 10}
 Best macro F1 (CV): 0.3582501478876265
 Decision Tree test accuracy: 0.36455293181215176
 Classification report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	538
1	0.33	0.37	0.35	538
2	0.20	0.19	0.20	538
3	0.21	0.27	0.24	539
4	0.18	0.16	0.17	539
5	0.22	0.17	0.19	538
6	0.40	0.39	0.39	539
accuracy			0.36	3769
macro avg	0.36	0.36	0.36	3769
weighted avg	0.36	0.36	0.36	3769



In []:

```
# Random Forest
# Random Forest with standardized preprocessing
import warnings
warnings.filterwarnings("ignore")

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, f1_score, classification_report
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.ensemble import RandomForestClassifier
from collections import Counter

# Step 1: Sample 100,000 rows (RF can handle larger sample sizes)
df_sample = df.sample(n=100000, random_state=42)

# Step 2: Collapse high-cardinality categorical variables
def collapse_categories(series, top_n=4):
    top_cats = series.value_counts().nlargest(top_n).index
    return series.where(series.isin(top_cats), other='Other')

for col in ['dIndustry', 'dOccup', 'dAncstry1', 'dAncstry2', 'iLang1']:
    if col in df_sample.columns:
        df_sample[col] = collapse_categories(df_sample[col], top_n=4)

# Cast other categorical predictors to string
categorical_vars_master = [
    "iSex", "dOccup", "dIndustry", "dAncstry1", "dAncstry2",
    "dHispanic", "dPOB", "iCitizen", "iEnglish", "iLang1",
    "iMilitary", "iSchool", "iWorklwk", "iDisabl1",
    "iDisabl2", "iMobillim", "iVietnam", "iWWII"
]
for col in categorical_vars_master:
    if col in df_sample.columns:
        df_sample[col] = df_sample[col].astype(str)

# Step 3: Target and predictors
y = df_sample['dTravtime'].astype('category')
predictors = [col for col in df_sample.columns if col not in y]
```

```
X = df_sample[predictors]
y_codes = y.cat.codes

# Step 4: Balance classes (undersample majority)
counts = Counter(y_codes)
min_count = min(counts.values())
balanced_idx = np.hstack([
    np.random.choice(np.where(y_codes == cls)[0], min_count,
        for cls in counts.keys()
])
X_balanced = X.iloc[balanced_idx]
y_balanced = y_codes.iloc[balanced_idx]

print("Balanced class counts:", Counter(y_balanced))

# Step 5: Train/test split
X_train, X_test, y_train, y_test = train_test_split(
    X_balanced, y_balanced, test_size=0.2, random_state=42,
)

# Step 6: Build preprocessing + RF pipeline
def make_pipeline_rf(predictor_list, n_estimators=200, max_depth=10):
    categorical_vars = [col for col in predictor_list if col in cat_cols]
    numeric_vars = [col for col in predictor_list if col not in cat_cols]

    preprocessor = ColumnTransformer(
        transformers=[
            ("num", StandardScaler(), numeric_vars),
            ("cat", OneHotEncoder(handle_unknown="ignore")),
        ]
    )

    pipeline = Pipeline(steps=[
        ("preprocess", preprocessor),
        ("model", RandomForestClassifier(
            n_estimators=n_estimators,
            max_depth=max_depth,
            random_state=42,
            n_jobs=-1
        ))
    ])
    return pipeline
```

```

# Step 7: Fit and evaluate
rf_pipeline = make_pipeline_rf(list(X_train.columns))
rf_pipeline.fit(X_train, y_train)
y_pred = rf_pipeline.predict(X_test)

acc = accuracy_score(y_test, y_pred)
macro_f1 = f1_score(y_test, y_pred, average="macro")
print("\nRandom Forest Results")
print("Test accuracy:", acc)
print("Macro F1:", macro_f1)
print("Classification report:")
print(classification_report(y_test, y_pred))

# Step 8: ROC curves
y_test_bin = label_binarize(y_test, classes=np.unique(y_test))
y_score = rf_pipeline.predict_proba(X_test)

fpr, tpr, roc_auc = {}, {}, {}
plt.figure(figsize=(8,6))
for i in range(y_test_bin.shape[1]):
    fpr[i], tpr[i], _ = roc_curve(y_test_bin[:, i], y_score[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])
    plt.plot(fpr[i], tpr[i], label=f"Class {i} (AUC = {roc_auc[i]:.2f})")

plt.plot([0,1],[0,1],'k--')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("Multiclass ROC (OvR) for Random Forest")
plt.legend()
plt.show()

# Step 9: Feature importance
model = rf_pipeline.named_steps["model"]
importances = model.feature_importances_
feature_names = rf_pipeline.named_steps["preprocess"].get_feature_names_out()

feat_imp = pd.DataFrame({"feature": feature_names, "importance": importances})
feat_imp = feat_imp.sort_values("importance", ascending=False)
print("\nTop 20 Feature Importances:")
print(feat_imp)

```

Balanced class counts: Counter({5: 2692, 0: 2692, 2: 2692, 4: 2692, 3: 2692, 1: 2692, 6: 2692})

Random Forest Results


Test accuracy: 0.3730432475457681

Macro F1: 0.3657846653920608

Classification report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	538
1	0.31	0.36	0.33	538
2	0.24	0.19	0.21	538
3	0.19	0.17	0.18	539
4	0.20	0.23	0.21	539
5	0.23	0.17	0.19	538
6	0.38	0.50	0.43	539
accuracy			0.37	3769
macro avg	0.36	0.37	0.37	3769
weighted avg	0.36	0.37	0.37	3769

Multiclass ROC (OvR) for Random Forest



Top 20 Feature Importances:

		feature	importance
3		num_dDepart	0.099622
21		num_iMeans	0.077229
7		num_dHours	0.054176
85	cat_dIndustry_0.0		0.039829
2		num_iClass	0.036025
122	cat_iWorklwk_1.0		0.035366
32		num_iRiders	0.032131
28		num_dRearning	0.031898
6		num_dHour89	0.028821
44		num_dWeek89	0.028194
46		num_iYearsch	0.024217
0		num_dAge	0.019907
18		num_iLooking	0.017821
9		num_dIncome1	0.017693
36		num_irPOB	0.015334
26		num_dPwgt1	0.015216
29		num_iRelat1	0.015089
35		num_dRpincome	0.014941
47		num_iYearwrk	0.014550
69	cat_iDisabl2_2		0.013041

In []: # Gradient Boosting

```

import warnings
warnings.filterwarnings("ignore")

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, f1_score, classification_report
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from xgboost import XGBClassifier
from collections import Counter

# Step 1: Sample 100,000 rows (XGB can handle larger sample)
df_sample = df.sample(n=100000, random_state=42)

# Step 2: Collapse high-cardinality categorical variables
def collapse_categories(series, top_n=4):
    if len(series.unique()) > top_n:
        series = series.value_counts()[:top_n].index
    return series

```

```

        top_cats = series.value_counts().nlargest(top_n).index
        return series.where(series.isin(top_cats), other='Other')

for col in ['dIndustry', 'dOccup', 'dAncstry1', 'dAncstry2', 'iL']:
    if col in df_sample.columns:
        df_sample[col] = collapse_categories(df_sample[col],)

# Cast other categorical predictors to string
categorical_vars_master = [
    "iSex", "dOccup", "dIndustry", "dAncstry1", "dAncstry2",
    "dHispanic", "dPOB", "iCitizen", "iEnglish", "iLang1",
    "iMilitary", "iSchool", "iWorklwk", "iDisabl1",
    "iDisabl2", "iMobillim", "iVietnam", "iWWII"
]
for col in categorical_vars_master:
    if col in df_sample.columns:
        df_sample[col] = df_sample[col].astype(str)

# Step 3: Target and predictors
y = df_sample['dTravtime'].astype('category')
predictors = [col for col in df_sample.columns if col not in y]
X = df_sample[predictors]
y_codes = y.cat.codes

# Step 4: Balance classes (undersample majority)
counts = Counter(y_codes)
min_count = min(counts.values())
balanced_idx = np.hstack([
    np.random.choice(np.where(y_codes == cls)[0], min_count,
                    replace=False) for cls in counts.keys()])
X_balanced = X.iloc[balanced_idx]
y_balanced = y_codes.iloc[balanced_idx]

print("Balanced class counts:", Counter(y_balanced))

# Step 5: Train/test split
X_train, X_test, y_train, y_test = train_test_split(
    X_balanced, y_balanced, test_size=0.2, random_state=42,
)

# Step 6: Build preprocessing + XGB pipeline
def make_pipeline_xgb(predictor_list, n_estimators=300, learn

```

```
categorical_vars = [col for col in predictor_list if col
numeric_vars = [col for col in predictor_list if col not

preprocessor = ColumnTransformer(
    transformers=[
        ("num", StandardScaler(), numeric_vars),
        ("cat", OneHotEncoder(handle_unknown="ignore")),
    ]
)

pipeline = Pipeline(steps=[
    ("preprocess", preprocessor),
    ("model", XGBClassifier(
        n_estimators=n_estimators,
        learning_rate=learning_rate,
        max_depth=max_depth,
        subsample=0.8,
        colsample_bytree=0.8,
        random_state=42,
        use_label_encoder=False,
        eval_metric="mlogloss"
    ))
])
return pipeline

# Step 7: Fit and evaluate
xgb_pipeline = make_pipeline_xgb(list(X_train.columns))
xgb_pipeline.fit(X_train, y_train)
y_pred = xgb_pipeline.predict(X_test)

acc = accuracy_score(y_test, y_pred)
macro_f1 = f1_score(y_test, y_pred, average="macro")
print("\nXGBoost Results")
print("Test accuracy:", acc)
print("Macro F1:", macro_f1)
print("Classification report:")
print(classification_report(y_test, y_pred))

# Step 8: ROC curves
y_test_bin = label_binarize(y_test, classes=np.unique(y_test)
y_score = xgb_pipeline.predict_proba(X_test)

fpr, tpr, roc_auc = {}, {}, {}
```

```
plt.figure(figsize=(8,6))
for i in range(y_test_bin.shape[1]):
    fpr[i], tpr[i], _ = roc_curve(y_test_bin[:, i], y_score[
        roc_auc[i] = auc(fpr[i], tpr[i])
        plt.plot(fpr[i], tpr[i], label=f"Class {i} (AUC = {roc_a

plt.plot([0,1],[0,1],'k--')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("Multiclass ROC (OvR) for XGBoost")
plt.legend()
plt.show()

# Step 9: Feature importance
model = xgb_pipeline.named_steps["model"]
importances = model.feature_importances_
feature_names = xgb_pipeline.named_steps["preprocess"].get_f

feat_imp = pd.DataFrame({"feature": feature_names, "importar
feat_imp = feat_imp.sort_values("importance", ascending=False)
print("\nTop 20 Feature Importances:")
print(feat_imp)
```

Balanced class counts: Counter({5: 2692, 0: 2692, 2: 2692, 4: 2692, 3: 2692, 1: 2692, 6: 2692})

XGBoost Results


Test accuracy: 0.3767577606792253

Macro F1: 0.37337428650665844

Classification report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	538
1	0.36	0.41	0.38	538
2	0.21	0.18	0.19	538
3	0.20	0.20	0.20	539
4	0.20	0.19	0.19	539
5	0.23	0.22	0.23	538
6	0.40	0.44	0.42	539
accuracy			0.38	3769
macro avg	0.37	0.38	0.37	3769
weighted avg	0.37	0.38	0.37	3769

Multiclass ROC (OvR) for XGBoost



Top 20 Feature Importances:

		feature	importance
85	cat_dIndustry_0.0	0.089014	
3	num_dDepart	0.059748	
7	num_dHours	0.031220	
59	cat_iCitizen_0	0.024493	
21	num_iMeans	0.024387	
32	num_iRiders	0.016737	
106	cat_dPOB_0.0	0.014239	
34	num_iRowndchld	0.012572	
11	num_dIncome3	0.009854	
28	num_dRearning	0.009527	
33	num_iRlabor	0.009399	
37	num_iRrelchld	0.009120	
9	num_dIncome1	0.008600	
35	num_dRpincome	0.008211	
115	cat_iSchool_2.0	0.008027	
27	num_iRagechld	0.008010	
40	num_iSept80	0.007999	
111	cat_dPOB_5.0	0.007953	
16	num_dIncome8	0.007888	
96	cat_iMilitary_3.0	0.007833	

```
In [ ]: # SVM with PCA
# Support Vector Machine (SVM) with standardized preprocessing
import warnings
warnings.filterwarnings("ignore")

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, f1_score, classification_report
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.decomposition import PCA
from sklearn.svm import SVC
from collections import Counter

# Step 1: Sample 50,000 rows (SVM is computationally heavier)
df_sample = df.sample(n=50000, random_state=42)
```

```

# Step 2: Collapse high-cardinality categorical variables
def collapse_categories(series, top_n=4):
    top_cats = series.value_counts().nlargest(top_n).index
    return series.where(series.isin(top_cats), other='Other')

for col in ['dIndustry', 'dOccup', 'dAncstry1', 'dAncstry2', 'iL']:
    if col in df_sample.columns:
        df_sample[col] = collapse_categories(df_sample[col], top_n=4)

# Cast other categorical predictors to string
categorical_vars_master = [
    "iSex", "dOccup", "dIndustry", "dAncstry1", "dAncstry2",
    "dHispanic", "dPOB", "iCitizen", "iEnglish", "iLang1",
    "iMilitary", "iSchool", "iWorklwk", "iDisab11",
    "iDisab12", "iMobillim", "iVietnam", "iWWII"
]
for col in categorical_vars_master:
    if col in df_sample.columns:
        df_sample[col] = df_sample[col].astype(str)

# Step 3: Target and predictors
y = df_sample['dTravtime'].astype('category')
predictors = [col for col in df_sample.columns if col not in y]
X = df_sample[predictors]
y_codes = y.cat.codes

# Step 4: Balance classes (undersample majority)
counts = Counter(y_codes)
min_count = min(counts.values())
balanced_idx = np.hstack([
    np.random.choice(np.where(y_codes == cls)[0], min_count,
                    replace=False) for cls in counts.keys()])
X_balanced = X.iloc[balanced_idx]
y_balanced = y_codes.iloc[balanced_idx]

print("Balanced class counts:", Counter(y_balanced))

# Step 5: Train/test split
X_train, X_test, y_train, y_test = train_test_split(
    X_balanced, y_balanced, test_size=0.2, random_state=42,
)

```

```

# Step 6: Build preprocessing + PCA + SVM pipeline
def make_pipeline_svm(predictor_list, C=1.0, gamma="scale"):
    categorical_vars = [col for col in predictor_list if col]
    numeric_vars = [col for col in predictor_list if col not

        preprocessor = ColumnTransformer(
            transformers=[
                ("num", StandardScaler(), numeric_vars),
                ("cat", OneHotEncoder(handle_unknown="ignore")),
            ]
        )

    pipeline = Pipeline(steps=[
        ("preprocess", preprocessor),
        ("pca", PCA(n_components=0.95, random_state=42)), #  

        ("model", SVC(kernel="rbf", C=C, gamma=gamma, proba=True))
    ])
    return pipeline

# Step 7: Fit and evaluate
svm_pipeline = make_pipeline_svm(list(X_train.columns))
svm_pipeline.fit(X_train, y_train)
y_pred = svm_pipeline.predict(X_test)

acc = accuracy_score(y_test, y_pred)
macro_f1 = f1_score(y_test, y_pred, average="macro")
print("\nSVM Results")
print("Test accuracy:", acc)
print("Macro F1:", macro_f1)
print("Classification report:")
print(classification_report(y_test, y_pred))

# Step 8: ROC curves
y_test_bin = label_binarize(y_test, classes=np.unique(y_test))
y_score = svm_pipeline.predict_proba(X_test)

fpr, tpr, roc_auc = {}, {}, {}
plt.figure(figsize=(8,6))
for i in range(y_test_bin.shape[1]):
    fpr[i], tpr[i], _ = roc_curve(y_test_bin[:, i], y_score[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])
    plt.plot(fpr[i], tpr[i], label=f"Class {i} (AUC = {roc_auc[i]:.2f})")

```

```
plt.plot([0,1],[0,1],'k--')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("Multiclass ROC (OvR) for SVM")
plt.legend()
plt.show()
```

Balanced class counts: Counter({5: 1331, 0: 1331, 2: 1331, 4: 1331, 3: 1331, 1: 1331, 6: 1331})

SVM Results


Test accuracy: 0.3717811158798283

Macro F1: 0.36535901736055854

Classification report:

	precision	recall	f1-score	support
0	0.99	0.99	0.99	266
1	0.34	0.37	0.35	266
2	0.20	0.21	0.20	266
3	0.22	0.17	0.19	267
4	0.22	0.23	0.23	266
5	0.23	0.16	0.19	266
6	0.35	0.47	0.40	267
accuracy			0.37	1864
macro avg	0.36	0.37	0.37	1864
weighted avg	0.36	0.37	0.37	1864

Multiclass ROC (OvR) for SVM



```
In [ ]: # Compare all models
from sklearn.metrics import accuracy_score, f1_score

def compare_models(X_train, X_test, y_train, y_test):
    results = []

    # Logistic Regression with backward stepwise selection
    acc_history = backward_stepwise(X_train, y_train, X_test)
    best_step = max(acc_history, key=lambda x: x[1])
    best_vars = best_step[2]
    best_pipeline = make_pipeline(best_vars)
    best_pipeline.fit(X_train[best_vars], y_train)
    y_pred = best_pipeline.predict(X_test[best_vars])
    results.append({
        "Model": "Multinomial Logistic Regression (Stepwise)",
        "Accuracy": accuracy_score(y_test, y_pred),
        "Macro F1": f1_score(y_test, y_pred, average="macro")
    })

    # KNN
    knn_pipeline = make_pipeline_knn(list(X_train.columns),
```

```
knn_pipeline.fit(X_train, y_train)
y_pred = knn_pipeline.predict(X_test)
results.append({
    "Model": "KNN (k=5)",
    "Accuracy": accuracy_score(y_test, y_pred),
    "Macro F1": f1_score(y_test, y_pred, average="macro"
})

# Random Forest
rf_pipeline = make_pipeline_rf(list(X_train.columns))
rf_pipeline.fit(X_train, y_train)
y_pred = rf_pipeline.predict(X_test)
results.append({
    "Model": "Random Forest",
    "Accuracy": accuracy_score(y_test, y_pred),
    "Macro F1": f1_score(y_test, y_pred, average="macro"
})

# XGBoost
xgb_pipeline = make_pipeline_xgb(list(X_train.columns))
xgb_pipeline.fit(X_train, y_train)
y_pred = xgb_pipeline.predict(X_test)
results.append({
    "Model": "XGBoost",
    "Accuracy": accuracy_score(y_test, y_pred),
    "Macro F1": f1_score(y_test, y_pred, average="macro"
})

# SVM
svm_pipeline = make_pipeline_svm(list(X_train.columns))
svm_pipeline.fit(X_train, y_train)
y_pred = svm_pipeline.predict(X_test)
results.append({
    "Model": "SVM (RBF)",
    "Accuracy": accuracy_score(y_test, y_pred),
    "Macro F1": f1_score(y_test, y_pred, average="macro"
})

# Convert to DataFrame for readability
results_df = pd.DataFrame(results)
return results_df
```

```
results_df = compare_models(X_train, X_test, y_train, y_test)
print(results_df)
```

Dropped dAge, accuracy 0.3540772532188841
 Dropped dAncstry1, accuracy 0.361587982832618
 Dropped iAvail, accuracy 0.3621244635193133
 Dropped iCitizen, accuracy 0.3621244635193133
 Dropped iDisabl1, accuracy 0.3621244635193133
 Dropped iFertil, accuracy 0.36319742489270385
 Dropped iClass, accuracy 0.3626609442060086
 Dropped iDisabl2, accuracy 0.36373390557939916
 Dropped iFeb55, accuracy 0.36373390557939916
 Dropped dHour89, accuracy 0.3642703862660944
 Dropped dHispanic, accuracy 0.3642703862660944
 Dropped dIncome4, accuracy 0.3642703862660944
 Dropped dIncome2, accuracy 0.36373390557939916
 Dropped dAncstry2, accuracy 0.36319742489270385
 Dropped iImmigr, accuracy 0.3626609442060086
 Dropped dIncome1, accuracy 0.3642703862660944
 Dropped iLang1, accuracy 0.36373390557939916
 Dropped iEnglish, accuracy 0.3648068669527897
 Dropped dIncome5, accuracy 0.3648068669527897
 Dropped iKorean, accuracy 0.365343347639485
 Dropped iLooking, accuracy 0.3648068669527897
 Dropped dIncome8, accuracy 0.3648068669527897
 Dropped iMay75880, accuracy 0.3648068669527897
 Dropped iMilitary, accuracy 0.365343347639485
 Dropped dIncome7, accuracy 0.3648068669527897
 Dropped iMarital, accuracy 0.3642703862660944
 Dropped iMobillim, accuracy 0.365343347639485
 Dropped iRagechld, accuracy 0.3648068669527897
 Dropped iMobility, accuracy 0.36587982832618027
 Dropped iRelat1, accuracy 0.36587982832618027
 Dropped dIncome6, accuracy 0.36587982832618027
 Dropped iPerscare, accuracy 0.3669527896995708
 Dropped iOthrser, accuracy 0.3669527896995708
 Dropped iRemplpar, accuracy 0.36641630901287553
 Dropped iRiders, accuracy 0.3696351931330472
 Dropped iRelat2, accuracy 0.3696351931330472
 Dropped iRlabor, accuracy 0.3696351931330472
 Dropped dRearning, accuracy 0.371781158798283
 Dropped iSchool, accuracy 0.37124463519313305
 Dropped iSubfam1, accuracy 0.3707081545064378

	Model	Accuracy	Macr
o F1			
o Multinomial Logistic Regression (Stepwise)	0.371781	0.36	

1939				
1	KNN (k=5)	0.315451	0.31	
6421				
2	Random Forest	0.368562	0.36	
3224				
3	XGBoost	0.354077	0.35	
1534				
4	SVM (RBF)	0.371781	0.36	
5359				